Why a merchant would choose us

Volume 2 — Product architecture. The commercial layer over the enterprise architecture: what we sell, how the pieces compose into products, and which of them a competitor cannot simply buy.

Document	EA-002 · v1.0 · Volume 2
Date	26 September 2026
Builds on	EA-001 v1.1 — Volume 1
Audience	Product, Commercial, Technology leadership
Products	10 lines, 38 sellable capabilities
Status	For review — 10 strategic decisions open

Contents

01 The thesis, in one page

02 The north star — five outcomes, one objective

03 Portfolio, capabilities, platform — and why they are three different counts

04 The product portfolio

05 Payment Optimization — the product the architecture was built for

06 Payment Intelligence — from reference data to a data moat

07 Credential Management and Credential Optimization

08 Commerce and Recurring — revenue recovery, not billing

09 Platform and Marketplace

10 Money — multi-currency, settlement and payout as products

11 Alternative and emerging rails, including stablecoin

12 Agentic commerce — what to build now and what to wait on

13 Developer and merchant experience

14 The intelligence loop

15 The compound moat

16 What a competitor can and cannot copy

17 The commercial tension nobody has priced yet

18 The blueprint — seven volumes, and what already exists

19 Ten strategic decisions

01

The thesis, in one page

> Volume 1 describes a platform that is well built. Well built is not a reason to switch provider. This volume is about the reason.

> A programmable payment operating system that separates policy, capability, strategy and execution; that continuously learns from payment outcomes; and that optimises authorization, authentication, credential selection, routing, cost and conversion on every transaction.

Three claims sit inside that sentence, and only the third is hard to copy.

CLAIM	WHAT IT MEANS	HOW HARD TO COPY
Programmable	Every payment decision is a versioned artefact rather than code in a release, so a change takes days.	Two years of platform work. Several competitors are already here.
Composable	Ten product lines built from four bounded capabilities on one shared platform, so a new method or rail is an addition rather than a rewrite.	Achievable with discipline. This is an architecture decision, not a moat.
Self-improving	Every transaction leaves counterfactual evidence, and the next decision is better because of it — measured against a reserved holdout.	Requires calendar time nobody can compress. This is the moat.

THE SENTENCE TO LEAD WITH

We do not just process your payment. We continuously improve how your payment gets processed, and we can prove by how much.

The second clause is the differentiated one. Plenty of platforms optimise. Very few can produce the number, the interval and the date — and a merchant who has been sold an optimisation story before will notice which half we are claiming.

What follows is the product layer that makes that claim true, and the four places where it is not true yet.

02

The north star — five outcomes, one objective

> A payment platform that optimises one thing will destroy another. Stating all five outcomes and the exchange rate between them is what turns optimisation from a slogan into a product.

ONE OBJECTIVE, FIVE OUTCOMES OPTIMISED TOGETHER

Authorization	Conversion	Cost	Risk	Experience
more transactions approved	**fewer buyers lost on the way**	**less paid per transaction**	**loss held inside appetite**	**less friction, more**
approval rate, first-attempt rate	checkout completion, authentication success	interchange, scheme fees, processor cost, FX	fraud rate, dispute rate, loss per unit	latency, step-up rate, settlement predictability

MERCHANT ECONOMICS

net margin per transaction, measured against a reserved holdout

— Optimising any one of the five alone is easy and usually destroys one of the other four. The platform exists to trade between them deliberately, with the exchange rate stated rather than implied.

Measured for three parties, not one

MERCHANT	US	CONSUMER
Authorization rate	Revenue per unit of volume	Success on first attempt
Checkout conversion	Gross margin after processor cost	Friction encountered
Cost per transaction	Processor and scheme economics	Authentication burden
Recovery rate on failures	Attach rate across product lines	Latency to a decision
Dispute and fraud loss	Retention and expansion	Predictability of outcome
Settlement predictability	Infrastructure and operating cost	Recourse when it goes wrong

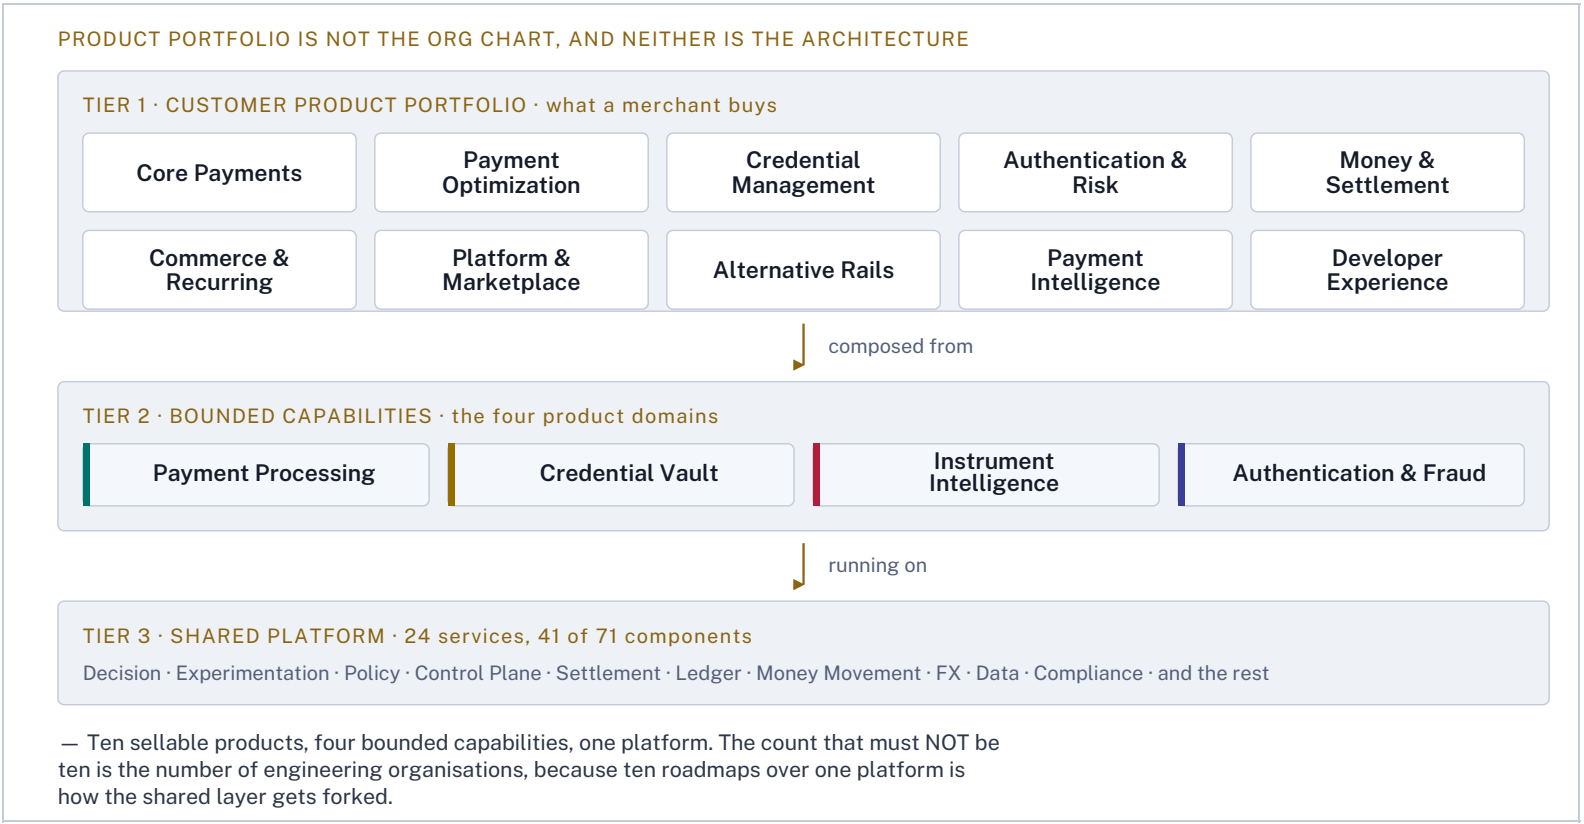
THE COLUMN THAT IS USUALLY MISSING

Most PSP optimisation narratives report the merchant column only. A change that raises approval by 37 bps and cost of acceptance by 2 bps is a clear merchant win, a marginal win for us on interchange-plus, and a loss for us on blended pricing.

A platform that cannot see all three columns at once will systematically ship changes that are good for the merchant and quietly bad for its own margin — and will not notice for several quarters. Section 17 is about this.

Portfolio, capabilities, platform — and why they are three different counts

> Ten products, four bounded capabilities, one platform. The number that must not become ten is the number of engineering organisations.



Volume 1 established that forty-one of seventy-one components are shared. That fact and a ten-product organisation are in direct tension: ten roadmaps competing for one platform team is the precise condition under which a shared service gets forked and the architecture quietly becomes ten architectures.

TIER	COUNT	WHAT IT IS	WHAT IT IS NOT
Product portfolio	10 lines, 38 capabilities	What a merchant buys, is entitled to and is billed for	Not a team structure and not a deployment boundary
Bounded capabilities	4 domains	Ownership and contract boundaries, from Volume 1	Not products — a merchant never buys “Instrument Intelligence”
Shared platform	24 services	The substrate every product line runs on	Not a cost centre — it is where the differentiation lives

THE ORGANISATIONAL RECOMMENDATION

Four to five product lines with named owners, not ten. Each line owns several sellable capabilities and negotiates with the platform as a customer. The platform organisation owns tiers two and three and publishes a roadmap the lines can plan against.

Ten product managers is fine and probably right. Ten engineering organisations with independent roadmaps is the failure mode, and the distinction has to be made explicitly because the org chart will otherwise follow the product catalogue by default.

04

The product portfolio

> Ten lines, thirty-eight sellable capabilities. Each maps to bounded capabilities from Volume 1 and to entries in the capability catalog — nothing here is sellable that the platform cannot evidence.

LINE	SELLABLE CAPABILITIES	BUILT FROM
Core Payments	Payment API · Order · Payment Intent · Attempt · Capture, refund, void · Payment methods	Payment Processing
Payment Optimization	Smart routing · Smart retry · Credential optimization · Authentication optimization · Cost optimization · Recovery	Decision Platform + all four domains
Credential Management	Vault · Network tokens · Account updater · PAR and identity · Consent and mandate	Credential Vault
Authentication & Risk	3DS and SCA · Exemptions · Delegated authentication · Fraud screening · Risk rules	Authentication & Fraud
Payment Intelligence	Instrument intelligence · Issuer intelligence · Cost intelligence · Benchmarking	Instrument Intelligence + Data Platform
Money	Settlement · Multi-currency · FX · Balances · Payouts · Ledger and reporting	Settlement, Sub-ledger, Money Movement, Currency
Commerce & Recurring	Subscriptions · Invoices · Instalments · Stored credentials · Revenue recovery	Payment Processing + Credential Vault + Decision Platform
Platform & Marketplace	Connected accounts · Split payments · Sub-merchant boarding · PayFac · Reserves and negative balance	Boarding, Risk, Money Movement, Tax, Compliance
Alternative Rails	ACH · RTP and FedNow · SEPA · Wallets · BNPL · Local methods · Stablecoin	Payment Processing + Connector Fabric + Money Movement
Developer & Merchant Experience	APIs and SDKs · Dashboard · Reporting and reconciliation · Sandbox · Self-service	Edge, Data Platform, Client Portal

THE RULE THAT KEEPS THE CATALOGUE HONEST

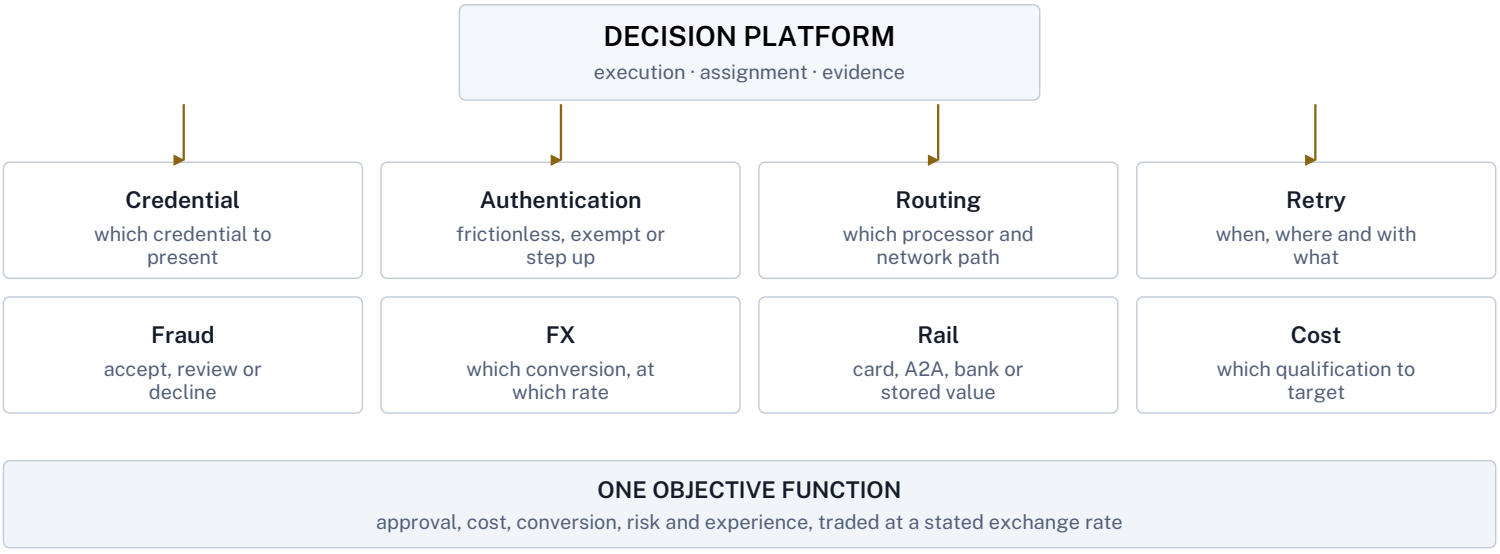
A capability becomes sellable only when the capability catalog says the platform can do it, verified rather than asserted. That is the three-registry separation from Volume 1 doing its job: engineering truth constrains commercial truth, and commercial truth constrains what a client can be entitled to.

The failure this prevents is the one every PSP has lived through — a capability sold in a region where it does not work, discovered at implementation.

Payment Optimization — the product the architecture was built for

> The Decision Platform is shared infrastructure. Payment Optimization is the product a merchant buys. Keeping those separate is what stops the platform becoming a feature of one product line.

ONE BRAIN, EIGHT DECISIONS, OPTIMISED JOINTLY



— Eight decisions optimised separately will fight each other: routing chases approval while cost optimisation undoes it, and retry recovers a transaction fraud would have declined. Joint optimisation is the whole point of one brain.

What is actually sold

SURFACE	WHAT IT OPTIMISES	THE MERCHANT-VISIBLE OUTCOME
Authorization	Processor and acquirer selection, cascade, soft-decline recovery, message enrichment, issuer-specific behaviour	Approval rate, first-attempt rate
Credential	PAN, network token, wallet token, merchant token, stored credential, account updater result	Approval rate and credential economics
Authentication	Frictionless 3DS, exemption selection, delegated authentication, step-up placement, authentication sequencing	Approval rate and checkout conversion
Retry and recovery	Timing, processor, credential form, rail, schedule	Recovered revenue
Cost	Interchange qualification, network fees, processor fees, FX cost, routing cost	Cost per transaction
Conversion	Checkout completion, authentication success, method presentation	Conversion rate

WHY THESE MUST BE ONE PRODUCT AND NOT SIX

Optimised separately they fight. Routing chases approval while cost optimisation undoes it. Retry recovers a transaction that fraud would have declined. Authentication reduces friction and raises chargebacks. Each is locally correct and the portfolio is worse.

One objective function with a stated exchange rate between approval, cost, conversion, risk and experience is the only way to trade between them deliberately. That exchange rate is a commercial input, owned by Product and Finance — it is decision 3 in Volume 1 and it blocks this entire product line.

Payment Intelligence — from reference data to a data moat

> Volume 1 treats instrument intelligence as reference data the platform reasons over. Commercially it is more than that: it is the one asset that gets better with volume and cannot be purchased.

The question the product answers: **what is this instrument, what can it do, where can it go, what will it cost, and how should we use it?** The first four are reference data and are purchasable. The fifth is learned from outcomes at our volume, and is not.

LAYER	FACTS	SOURCE	COPYABLE?
Identity	BIN, range, issuer, product, country, currency, network	Scheme and vendor data	Yes — anyone can buy it
Capability	Token eligibility, 3DS capability, authentication support, routing eligibility	Verified with processors and networks	Partly — verification takes effort
Economics	Interchange qualification, scheme fees, processor cost, FX exposure	Scheme publications plus our own rates	Partly
Regulatory	Regulated debit status, cross-border flags, scheme classifications	Published rules	Yes
Behaviour	**Observed issuer approval patterns, authentication posture, stand-in behaviour, decline reason semantics, recovery response**	**Our own transaction outcomes**	**No — requires volume and time**

THE LINE BETWEEN A DATASET AND A MOAT

The first four layers are a procurement exercise. The fifth is a function of how many transactions we see, over how long, with outcomes joined at four horizons and a reserved holdout to measure against.

That is why the evidence requirements in Volume 1 are not an engineering nicety. Without propensity and the fallback counterfactual recorded at decision time, the behaviour layer is observational data about what we happened to do — not evidence about what would have happened otherwise. The first can describe; only the second can optimise.

07

Credential Management and Credential Optimization

> Two products from one bounded capability. One is infrastructure a merchant needs; the other is an outcome a merchant pays for.

CREDENTIAL MANAGEMENT	CREDENTIAL OPTIMIZATION
Vault — surrogate and credential state	Which credential form to present, per transaction
Network tokens, wallet tokens, MPAN	Token versus PAN, decided on measured outcome rather than policy
Account updater and lifecycle	When to refresh, and whether to retry on the refreshed credential
PAR and cross-channel identity	Which identity anchor to use per channel
Consent and mandate	Which stored-credential framing maximises approval within the mandate
Sold as capability	*Sold as outcome, measured against a holdout*

The split matters commercially because credential management is table stakes — every serious PSP has a vault — while credential optimization is a claim almost nobody can evidence. The first experiment in the programme exists precisely to produce that evidence for our own vault, and the number it produces is the number the renewal conversation needs.

THE OMNICHANNEL CONSTRAINT

A wallet DPAN and its funding FPAN derive different surrogates; only PAR links them. Cross-channel identity built on the surrogate alone treats every in-store wallet tap as a new customer.

This is decision 4 in Volume 1 — the surrogate tweak scope — and it is a one-way door. Per-merchant scoping gives privacy and forecloses cross-merchant omnichannel identity. Per-platform scoping does the reverse. Changing it later re-tokenises the estate.

Commerce and Recurring — revenue recovery, not billing

> A billing API is a commodity. Recovering the revenue that recurring billing loses is not, and it is the same optimisation machinery pointed at a different problem.

RESOURCE	WHAT IT ADDS	WHY IT IS NOT JUST A PAYMENT
Customer	The durable relationship a subscription hangs from	Outlives any credential
Subscription	Term, plan, status, lifecycle	Has a state machine payments do not
Plan	Price, interval, trial, proration	Pricing logic distinct from transaction pricing
Invoice	What is owed for a period, before it is collected	Can exist unpaid, be partially paid, be credited
Payment Intent	The collection attempt for an invoice	Already in the Volume 1 model
Retry schedule	When and how to try again	The optimisation surface

Where the optimisation lives

Recurring failure is rarely a single decision. It is a sequence: which credential to present, whether to refresh it first, when to retry, on which processor, with which authentication framing, and when to stop. Every one of those is a decision point, which means every one is a strategy, which means every one is experimentable.

DECISION	STRATEGY EXAMPLES	MEASURED AS
Credential form	Network token first, PAN on refresh failure, updated credential when available	Recovery rate
Refresh timing	Pre-emptive updater call before a known expiry versus reactive on decline	Recovery rate and updater cost
Retry schedule	Fixed ladder versus issuer-specific timing versus learned	Recovery rate per attempt
Processor	Same processor versus alternate on soft decline	Recovery rate
Authentication	Whether to authenticate a retry, and under which exemption	Recovery rate and dispute exposure
Stop rule	When further attempts destroy more value than they recover	Cost per recovered unit

THE PRODUCT THIS BECOMES

Not “subscription billing”. ***Revenue recovery***, sold on a measured recovery rate against a holdout, with the failed-payment population as the denominator.

That framing also changes the pricing conversation: a share of recovered revenue is defensible in a way that a per-invoice fee is not, and it aligns our incentive with the merchant’s. It is also only sellable if the measurement is real, which brings it back to the holdout.

09

Platform and Marketplace

> The segment with the highest growth and the longest list of prerequisites. Eight of the nine supporting capabilities are the additions from Volume 1, which is why they were not optional.

THE PLATFORM MODEL, AND WHAT EACH SHAPE NEEDS

PLATFORM — a client whose own clients take payments · marketplace, ISV, franchise, PayFac

Connected account the sub-entity that receives funds	**Collect** money taken on the platform's behalf or the seller's	**Split** allocation across sellers, platform fee and tax
Hold & reserve risk-driven retention against future loss	**Net & offset** across sellers, refunds and negative balances	**Payout** to each entity, on its own schedule and currency

AND WHAT EVERY ONE OF THEM DRAGS IN

KYC / KYB / UBO	Underwriting	Merchant hierarchy	Balances & ledger
Negative balance	Fee & split rules	Tax & 1099	Dispute allocation

— A platform proposition is not a payments feature. Eight of these nine supporting capabilities are the additions from Volume 1 section 15, which is why they were not optional.

SHAPE	WHO TAKES THE MONEY	WHO CARRIES THE RISK	WHAT WE MUST PROVIDE
Marketplace	Platform collects, then distributes	Platform, usually	Split, hold, net, payout, tax, dispute allocation
PayFac	Platform is the merchant of record for its sellers	Platform	Sub-merchant boarding, underwriting, reserves, 1099
ISV / referral	Seller collects directly	Seller	Boarding, revenue share, reporting
Franchise	Central or local depending on the model	Mixed	Hierarchy, consolidated and per-outlet reporting
Connected accounts	Either, per transaction	Configurable	Balances, negative balance handling, per-entity payout

THE UNGLAMOROUS BLOCKER

Merchant hierarchy. Every one of split, net, reserve, payout, tax, dispute allocation and reporting joins on it. Volume 1 decision 8 asks who owns it, and the answer has to be Reference and Master Data rather than Boarding — Boarding is the first writer, not the owner, and hierarchy changes long after boarding through acquisitions, restructures and new outlets.

This is the most common master-data failure in acquiring and it surfaces at month end, in finance, after the platform proposition has already been sold.

Money — multi-currency, settlement and payout as products

> Volume 1 models seven currency contexts and the conversions between them. Commercially, each conversion is a moment where margin is either captured or leaked, and each is a decision point.

PRODUCT	WHAT THE MERCHANT GETS	WHERE THE OPTIMISATION IS
Multi-currency acceptance	Charge the buyer in their currency	Presentment currency selection and its effect on conversion
Multi-currency processing	Authorise in the currency that performs best	Processing currency selection — approval and cost differ by acquirer and corridor
Multi-currency settlement	Settle in chosen currencies, not forced conversion	Settlement currency selection and timing
FX	A rate with a stated markup, not an opaque spread	Rate source selection, quote duration, who bears the spread
DCC	Offer the buyer their own currency at the point of sale	Whether to offer, and to whom — a conversion and disclosure trade
Balances and payout	Hold, net and pay out per entity and currency	Payout timing and currency, netted against exposure

FX AS AN OPTIMISATION SURFACE

Once the seven-currency chain is explicit, which conversion happens where, at whose rate, with what markup, and who bears the spread becomes a strategy bound at a decision point rather than a configuration setting.

That is a genuine differentiator and it is unavailable to any platform that models currency as an attribute of an amount. It is also where a meaningful share of PSP margin sits, which is why it needs a named owner rather than a default.

Alternative and emerging rails, including stablecoin

> The architectural position is right and worth stating plainly: a payment is an intent, and a rail is an execution choice. Stablecoin is another rail. The operational position is more complicated.

RAIL	FINALITY	REVERSIBILITY	WHAT IT BREAKS IN OUR MODEL
Card	Deferred — clearing and settlement	Chargeback for months	Nothing. The model assumes it.
ACH	Deferred, days	Return windows by code	Settlement timing and exposure modelling
RTP / FedNow	Immediate	None — request for return only	Reserve logic assumes a claw-back path
SEPA	Deferred or instant by scheme	Mandate-dependent	Mandate management across schemes
Wallet	Underlying rail	Underlying rail	Nothing new
BNPL	Provider-dependent	Provider-dependent	Dispute allocation with a third party
Stablecoin	**Immediate and irreversible**	**None**	**Dispute, reserve, recon and refund all assume reversibility**

"JUST ANOTHER RAIL" IS ARCHITECTURALLY RIGHT AND OPERATIONALLY MISLEADING

The connector fabric, money movement and settlement abstractions do accommodate stablecoin without redesign. That part of the argument holds, and building a separate stablecoin stack would be a mistake.

What does not carry over is everything downstream of irreversibility. A refund becomes a new outbound payment rather than a reversal. A dispute has no claw-back path, so risk moves entirely to pre-authorisation. Reserves cannot be recovered from a settled transaction. Reconciliation is against a chain rather than a clearing file. And the compliance surface — travel rule, wallet screening, custody — is different in kind, not degree.

So: one rail abstraction, yes. One set of operational assumptions, no. That distinction should be made before the first stablecoin transaction, not after the first irreversible mistake.

Agentic commerce — what to build now and what to wait on

> The shift from “a person clicks buy” to “a person authorises an agent to buy within constraints” changes what an authorization has to carry. Some of that is schema and should be decided now; most of it is product and should not.

THE TRUST CHAIN AN AGENT PURCHASE HAS TO CARRY

Human	the principal, who holds the money and the liability
Delegation	what authority was granted, to whom, for how long, with what limits
Agent	the acting party — identified, attested, revocable
Commerce intent	what is being bought, from whom, at what price
Payment intent	the amount, carrying the mandate and the limits
Credential	a token scoped to the mandate, not the cardholder’s full credential
Authorization	with evidence of delegated authority attached to the message
Disputability	who is liable when the agent was wrong, and how it is proven

— Three of these are schema decisions: delegation, mandate-scoped credential, and authorization evidence. Those are cheap now and very expensive to retrofit. The rest is product, and the standards are still moving.

The split

DECIDE NOW — schema	WHY IT CANNOT WAIT
Delegation record — who granted what authority, to whom, for how long, with what limits	It is a first-class resource. Bolting it on later means every existing mandate lacks it.
Mandate-scoped credential — a token bound to the delegation, not the full cardholder credential	Credential scoping is a vault schema decision and a token lifecycle decision. Retrofitting is re-provisioning.
Authorization evidence — proof of delegated authority carried with the message	If it is not captured at authorization it cannot be reconstructed, exactly like propensity.
Agent identity as a distinct party from merchant and consumer	Every downstream model — risk, dispute, reporting — assumes two parties. Adding a third later touches all of them.

WAIT — product	WHY
Agent-facing APIs and developer surface	The standards are moving. Building a product surface on a moving standard buys a rewrite.
Agent-specific risk models	No outcome data exists yet. A model with no evidence is a rule with a confidence interval attached.
Agent dispute and liability rules	Scheme and regulatory positions are not settled. Ours should not be either.
Agent commerce as a sellable line	Until the trust model is settled we would be selling an integration, not an outcome.

THE SAME ARGUMENT AS PROPENSITY LOGGING

Four schema decisions are cheap today and very expensive to retrofit. Everything else is a product bet on a standard that is still forming, and the current state of those standards needs a fresh scan before any of it is committed — this document should not be the source of truth on where the protocols have landed.

The recommendation is therefore narrow and unglamorous: build the four primitives into the resource model now, ship nothing agent-facing, and revisit the product question with a current view of the standards.

Developer and merchant experience

> The dimension that decides whether any of the above is ever reached. A platform that takes six weeks to integrate loses to one that takes six hours, whatever its approval rate.

SURFACE	THE BAR	WHY IT IS AN ARCHITECTURE QUESTION
API	One resource model across every product line, consistent pagination, errors, idempotency and versioning	Ten product lines with ten API styles is the symptom of ten engineering organisations
SDKs	First-class in the languages the segment actually uses, generated from one specification	Generated SDKs require one specification, which requires one resource model
Sandbox	Deterministic simulators for processor, network and issuer behaviour, including failures	This is the Sandbox and Simulator Fabric from Volume 1 — it is infrastructure, not documentation
Documentation	Task-shaped rather than endpoint-shaped, with runnable examples	Cannot be written coherently over an inconsistent API
Dashboard	Transaction search to full decision trace in under two minutes, self-serve	This is the decision trace requirement from Volume 1, exposed to the client
Reporting and recon	Merchant-level statements that reconcile to the sub-ledger without support	Requires the settlement, ledger and money movement separation to be real
Self-service	Config, users, keys, webhooks, limits and product activation without a ticket	Requires entitlement to be a grant store rather than scattered configuration

THE DECISION TRACE IS A PRODUCT FEATURE, NOT JUST A SUPPORT TOOL

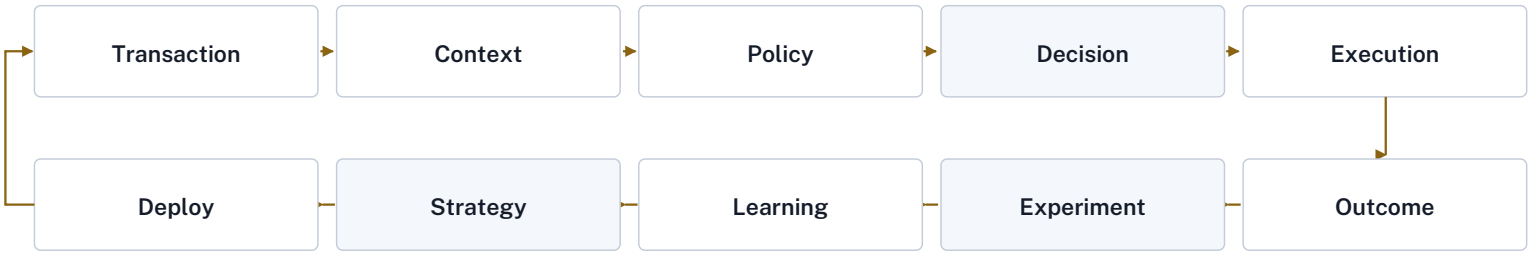
Volume 1 specifies a per-transaction decision trace: membership, arm, flow, every decision applied and what the baseline would have decided. Built for internal explainability, it is also the single most differentiating thing a merchant-facing dashboard can show.

“Why was this transaction declined” is the most common merchant support question in acquiring, and almost nobody can answer it without an engineer. Exposing the trace turns an internal requirement into a visible product advantage at close to zero marginal cost.

The intelligence loop

> Volume 1 has the parts. This is the claim they add up to, and it should be the headline rather than the appendix.

THE PSP IMPROVES EVERY TIME IT PROCESSES A PAYMENT



WHAT MAKES THE LOOP CLOSE RATHER THAN JUST TURN

Propensity on every decision
without it, nothing downstream can be evaluated off-policy

Fallback counterfactual
what the baseline would have done, recorded at decision time

A reserved holdout
the only way to measure the platform rather than one experiment

One epoch per decision
so a decision is reproducible months later

— None of the four can be added retrospectively. A loop missing any one of them is a reporting pipeline that looks like a learning system.

> The platform improves every time it processes a payment.

That is a materially stronger statement than “our platform uses AI”, and unlike that claim it is falsifiable. The holdout is what makes it falsifiable, which is why reserving it is the one decision Volume 1 asks for today.

WHY THE LOOP IS THE STRATEGY AND NOT A FEATURE

Every other advantage in this document depreciates. Connectivity gets matched. Credential capabilities get matched. A routing algorithm gets matched within a release cycle of someone noticing.

An evidence base accumulated under a reserved holdout appreciates, and it cannot be bought or compressed — it takes the calendar time it takes. A competitor starting today is three months behind before they can evaluate anything at all, and that gap widens rather than closes.

The compound moat

SIX MOATS THAT COMPOUND INTO ONE

Data
instrument intelligence joined to transaction outcomes, at our volume

Decision
one place where every payment decision is made, versioned and measured

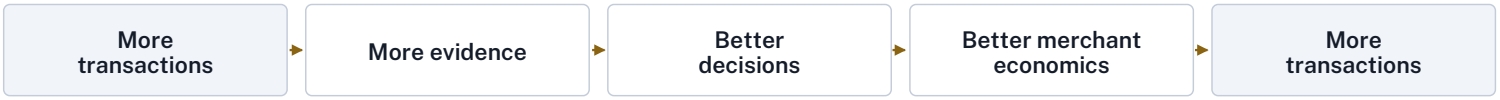
Credential
vault, tokens, lifecycle and cross-channel identity in one boundary

Network
processor, acquirer and scheme connectivity, and the capability truth about it

Economics
interchange, processor cost, FX and routing modelled together rather than separately

Learning
every transaction leaves evidence the next decision uses

AND THE FLYWHEEL THEY DRIVE



— Only the last two moats are hard to copy. Connectivity and credentials can be bought; evidence accumulated under a reserved holdout cannot, because it takes calendar time nobody can compress.

MOAT	SOURCE	TIME TO COPY WITH FUNDING
Network	Processor, acquirer and scheme connectivity	12–24 months, and buyable
Credential	Vault, tokens, lifecycle, identity	12–18 months, and buyable
Decision	One decision platform, versioned and measured	18–30 months, hard to buy
Economics	Interchange, processor cost, FX and routing modelled together	Depends on scale, partly buyable
Data	Instrument intelligence joined to outcomes at our volume	Requires our volume
Learning	**Evidence under a reserved holdout**	**Requires calendar time no amount of funding compresses**

What a competitor can and cannot copy

> Worth being honest internally about which parts of this story are genuinely defensible, because overclaiming the copyable parts is how a differentiation narrative loses credibility with the people who have to build it.

CLAIM	HONEST ASSESSMENT
“We have smart routing”	Everyone says this. Several competitors do it well. Not a differentiator.
“We have a vault and network tokens”	Table stakes. Absence is a disqualifier; presence is not an advantage.
“We support many payment methods”	A procurement and integration exercise. Copyable with money.
“We use AI”	Unfalsifiable, and increasingly heard as noise. Avoid.
“We can change a payment decision in days”	True and meaningful, but reachable by anyone who builds the control plane.
“We can prove a configuration was better”	**Few can. Requires propensity, counterfactual, holdout and the discipline to keep them.**
“We know how this issuer behaves for this instrument, measured”	**Requires volume and accumulated evidence. Not purchasable.**
“Here is what we were worth to you last quarter, against a holdout”	**Almost nobody can say this. It is the strongest sentence available to us.**

The commercial tension nobody has priced yet

> The optimisation story has an economic problem inside it, and it is better to surface it now than to discover it in a margin review.

Optimisation that raises approval frequently raises cost of acceptance too. Whether that is good for us depends entirely on the pricing model:

PRICING MODEL	MERCHANT OUTCOME	OUR OUTCOME	INCENTIVE ALIGNMENT
Interchange-plus	Gains the approval, pays the added cost	Gains revenue on incremental volume	Aligned
Blended	Gains the approval, pays nothing extra	Absorbs the added cost	**Opposed**
Optimisation as a premium SKU	Pays for optimisation explicitly	Revenue from the SKU	Aligned, but creates a reason not to optimise the base product
Share of measured gain	Pays a share of what was gained	Revenue proportional to value created	Aligned, and requires the measurement to be real

THREE THINGS THAT FOLLOW

One. On blended pricing, a successful optimisation can be a net loss to us. Worked through in the credential experiment: roughly $1.16M a year to the merchant and about $129K of cost to us per $1B of qualifying volume. Unpriced, that scales with success.

Two. Selling optimisation as a premium tier creates a perverse incentive to keep the base product unoptimised. That is commercially rational and strategically corrosive, and it should be named before someone proposes it as a packaging idea.

Three. Share-of-gain pricing is the only model that aligns all three parties, and it is only possible if the gain is measured against a holdout. The measurement discipline is therefore not just an engineering requirement — it is what makes the best pricing model available at all.

The blueprint — seven volumes, and what already exists

> Two of these are written and three exist in substantial draft. The gap is smaller than a seven-volume plan implies.

VOLUME	SCOPE	STATUS
1 — Enterprise Architecture	Components, boundaries, rules, runtime, control plane	**Complete** — EA-001 v1.1
2 — Product Architecture	Portfolio, product lines, packaging, entitlement, differentiation	**This document** — EA-002 v1.0
3 — Decision Platform	Decision graph, strategies, experimentation, counterfactuals, features, explainability	Substantially drafted — experimentation and decision platform PRD, routing strategy PRD, policy platform PRD
4 — Global Money	Settlement, ledger, FX, balances, payouts, rails	Partial — covered architecturally in Volume 1, not yet as a design
5 — Commerce Platform	Orders, intents, attempts, recurring, marketplace, platform payments	Partial — resource model work exists; recurring and marketplace do not
6 — Merchant Experience	Onboarding, developer experience, dashboard, reporting, self-service	**Gap** — the least developed area
7 — Operating Model	Product and platform organisation, governance, experimentation governance, commercialisation	Substantially drafted — platform governance and operating model

WHAT I WOULD WRITE NEXT, AND WHY

Volume 6. It is the emptiest and it gates adoption of everything else. A merchant never experiences the decision platform; they experience the integration, the dashboard and the statement. It is also the volume where our competitors are strongest, which makes it the most uncomfortable one to leave last.

Then Volume 4. Money is architecturally covered and commercially unspecified, and it is the volume that blocks the platform and marketplace proposition.

Ten strategic decisions

> Deferred here from Volume 1. Each is a product or commercial decision that constrains the architecture and is constrained by it. None resolves itself.

#	DECISION	BLOCKS	RECOMMENDATION
10	Payment Intent versus Order semantics — who owns which	Every API surface	Order is commerce, Intent is collection. Payment Processing owns both, Commerce composes them.
11	The canonical payment method abstraction	Time to add a rail	Method is the client's choice; rail is ours. Never expose the rail as an API branch.
12	The canonical multi-currency model	FX margin	The seven-currency chain from Volume 1, with each conversion a decision point.
13	Where recurring billing lives	Whether revenue recovery is a product	A composition over Payment Processing and Credential Vault, not a fifth domain.
14	Where marketplace and split orchestration lives	The platform proposition	A product line composing Money Movement, Risk, Tax and Boarding — not a new bounded capability.
15	The global acquiring model — own licences, sponsored, or hybrid	Market entry sequencing	Hybrid by market. Needs a market-by-market plan, not a principle.
16	The SLO for platform-added authorization latency	Whether the budget is real	Adopt the 68 ms budget in Volume 1 as a target and 100 ms as a ceiling, with per-stage owners.
17	The Decision Platform feature and model architecture	Learned strategies	Volume 3. Blocked on nothing except starting.
18	How counterfactual evaluation works and who validates it	Whether a model reaches production safely	Volume 3. Needs an owner outside the team that builds the models.
19	The agentic commerce trust model	Nothing today; everything in two years	Build the four schema primitives now. Ship no agent-facing product until the standards settle.

THE ONE TO SETTLE FIRST

Decision 16. Not because latency is the most strategic question here, but because it is the only one on this list that silently becomes unfixable. Every product line in section 4 adds work to the authorization path, and without a ceiling with named per-stage owners the budget is consumed by whoever ships first.

The others can be argued about for a quarter without permanent cost. This one cannot.